

A parallelised and distributed approach to implementing Conway's Game of Life

AARON CHAN

University of Bristol
ho21739@bristol.ac.uk

FERDINAND HUBBARD

University of Bristol
ej21378@bristol.ac.uk

December 5, 2022

Abstract

We intend to explain our two solutions for implementing Conway's Game of Life in Golang, namely the Parallel and Distributed versions. Within this report, we will discuss the impacts of thread usage in both our implementations, as well as the effect of distributed computation as opposed to single-machine computation. In essence, we found that execution time was sped up as we increased thread count, reduced branch instructions and implemented a halo exchange mechanism.

I. INTRODUCTION TO THE PARALLEL SOLUTION

As an initial attempt, we began by implementing a single-threaded Game of Life (GoL) Engine, expanding upon this work with parallelisation by virtue of concurrent go-routine workers. Recognising that there exists various methods to calculating the perimeter for `getNeighbourCount`, we tried various counting methods - taking into account the modularity of Game of Life - as well as finding channel alternatives.

II. BENCHMARKING METHODOLOGY

Our means of evaluating parallelisation performance gains was via the Go testing framework. Using 512px x 512px PGM images, we repeated each benchmark 10 times, taking the average runtime for each thread to acquire the central value, and plotting them. Each feature implementation was benchmarked in the same fashion, using Linux lab machines with CPUs being Intel i7-8700 (6 core, 12 threads).

III. A SPLIT PERSONALITY PROGRAM

A. Parallelisation logic

The baseline implementation's distributor function starts with reading the image into 2D slice via `io` commands, before evenly dividing the workload between workers. With the aid of go-routines, workers have




Figure 1: Modulo Implementation

access to the old slice as well as the boundaries in which to process. By using `getNeighbourCount`, we apply the Game of Life laws onto the world, piping results into a blank slice of size:

$$(endRow - startRow) \times ImageWidth$$

Sending the results back to the distributor, we integrate back together the new slices, move the pointer to the old world to this newly-merged world and send the appropriate signals on the `events`, `io` channels.

B. On the topic of analysis 1

From Fig. 1, we can conclude that multithreading have an overall improvement on performance. There is a correlation between increasing number of threads and decreasing execution times, likely due to some functions, such as `getNeighbourCount`,

being "embarrassingly parallelisable". We favour the term "delightfully parallel". As we have more cores for processing, the time taken for each turn is reduced, hence the curve.

We do take notice of the obvious trend of diminishing returns - the gradient of the time-thread curve flattens. We attribute this to unparallelisable core tasks such as IO - the major slowdown that is reading and writing to/from the disk, and submit that using more threads than a CPU physically has causes overhead in scheduling, minimising the benefits of allocating extra threads for processing. We suspect also that, contributing to this, the workload division does not change as much the more threads there are, which explains the reciprocal graph shape curve.

Regarding the numerical results, we find greatest performance was with 16 threads (despite the diminishing returns) with a 67.6% decrease in processing time. As for the worst change in performance thread-to-thread, we see the 14th - 15th threads having the smallest difference in performance change (0.353%), but acknowledge the potential anomalous data point.

IV. HOW TO COUNT?

A. Counting methods

It is well documented that division and other non-atomic arithmetic operations are slow. Thus, we submit our alternative implementations for `getNeighbourCount`. Fig. 1, our original implementation, uses the modulo operation (%), whereas Fig. 2, Fig. 3, use Branching (if-statements) and Bitmasking respectively. From this data, we see that Bitmasking is the fastest on average, but we can identify positive notes from all three.

B. On the topic of Analysis 2

Statistically-speaking, we find the solution using modulo operations is most affected by thread usage, having standard deviation 413ms, as opposed to 107ms and 73ms for branching and bitmasking respectively. Interested persons will find our results achieved by:

$$SD(\forall x \in \text{runtimes} : \max(\text{runtimes}) - x)$$

As discussed before, this, however, implies not that higher threads means greater performance as there exists other, non-parallelisable tasks setting a base level of execution time. The low standard deviation




Figure 2: Branching Implementation




Figure 3: Bitmask Implementation

of the bitmasked implementation alludes to this portion of the program being as near optimised as we possibly can.

When it comes to average times for threads, bitmasking takes the lead with 890ms, followed by 1026ms and 1989ms for branching and modular respectively. This does identify branching for the greatest decrease in runtime, with a 74.2% decrease, as opposed to 32.1%.

C. Our Algorithms

Code-wise, our initial implementation for `getNeighbourCount` used modulo operators unnecessarily. It did the following:

1. Set row = coordinate row % ImageHeight
2. Set col = coordinate col % ImageWidth
3. If row | col < 0, then add ImageHeight/Width respectively
4. If world[row][col] was 255, increment alive count

With this, while it succeeds in its purpose, it does not succeed efficiently. Areas of improvement include removal of the modulo operator, which, here,

does not serve any real purpose. This leads to our second implementation - branching:

Algorithm 1 Branching getNeighbourCount

```

Let  $i, j \leftarrow (row, column)$  on world
Let  $O \leftarrow [-1, 0, 1] \times [-1, 0, 1] - \{0, 0\}$ 
Let  $alive \leftarrow 0$ 
for  $(x, y) \in O$  do
   $row \leftarrow (i + y)$ 
  if  $row < 0$  then
     $row \leftarrow ImageHeight - 1$ 
  else if  $row = ImageHeight$  then
     $row \leftarrow 0$ 
  end if
   $col \leftarrow (j + x)$ 
  if  $col < 0$  then
     $col \leftarrow ImageWidth - 1$ 
  else if  $col = ImageWidth$  then
     $col \leftarrow 0$ 
  end if
  if  $world[row][col] = 255$  then
     $alive++$ 
  end if
end for
return  $alive$ 

```

Although the line count has certainly increased, the performance has drastically improved. Citing the 74.2% decrease, this does appear to be "good enough" as an implementation. An improvement may be to rid ourselves of branches, as this involves branch prediction which itself is intensive, due to how the CPU does not know which path will be taken ahead of time. If we use only logical operations, we get the Bitmasking implementation:

Algorithm 2 Bitmasked getNeighbourCount

```

Let  $i, j \leftarrow (row, column)$  on world
Let  $O \leftarrow [-1, 0, 1] \times [-1, 0, 1] - \{0, 0\}$ 
Let  $alive \leftarrow 0$ 
for  $(x, y) \in O$  do
   $row \leftarrow (i + y) \& (ImageHeight - 1)$ 
   $col \leftarrow (j + x) \& (ImageWidth - 1)$ 
   $alive \leftarrow alive + world[row][col] >> 7$ 
end for
return  $alive$ 

```

This algorithm is our fastest implementation. Minimising both line count and runtime, it counts the neighbours using purely bitwise operators and assignments - no branching needed.




Figure 4: Memory Sharing

V. MEMORY SHARING

A. Motivation

Looking at our achievements in counting, we noticed overhead associated with channel communication. Exploring Go's implementation for channels, we see much overhead. Since we know a channel is a thread-safe memory sharing mechanism, we can rid the overhead, using only mutex locks and condition variables.

B. Implementation

Our homemade channel implementation does the following:

1. Acquire the mutex lock
2. While `condition` is false, await broadcast unless non-blocking send
3. Use the value, then broadcast to other processes
4. Finally, unlock mutex

We introduce 2 methods, `Send` & `Receive`. `Send` waits for the memory space to be empty, and replaces it with a value. `Receive` waits for the memory space to point to a value, and then replaces it with `nil`. Go 1.12 does not have generics, which vastly complicated implementation. Benchmarks with memory sharing (and bitmasked `getNeighbourCount`) is shown in Fig. 4.

C. On the topic of Analysis 3

This data was not very analysable. The lack of clear trend and sporadic data points suggests no real difference in performance when adding memory sharing, or faulty implementation (which we realistically do consider a possibility). A trend, if any, we identify

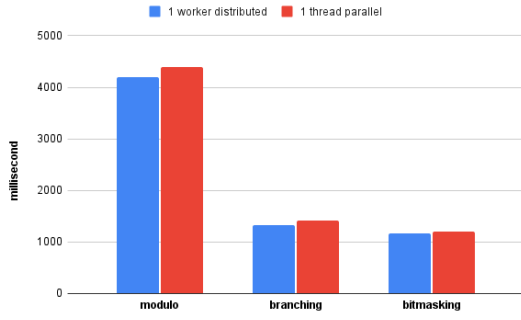


Figure 5: 1 worker vs. 1 thread

is that the Modulo performance suffered, but due to the lack of clarity, this is inconclusive. We add that repeated trials yielded similar style results. Our method used the following:

$$\forall t \in T : ORIG[t] - MEM[t]$$

Where *MEM* is memory sharing implementation time, *ORIG* is channel impl. time, *T* is the number of threads.

VI. DISTRIBUTED IMPLEMENTATION

Similar to how we did the parallel side, we started the distributed section by using a single worker, this time with the worker being a separate process and communicated to via Remote Procedure Calls (RPC henceforth). This expanded to doing using multiple workers, implementing fault tolerance and finally halo exchange.

VII. A LONESOME WORKER

A. Logic

Our single-worker distributed implementation moved much of the processing to said worker. From the distributor, we sent a single RPC call to the worker, processed GoL turns on the worker, then sent the final state back. Results are shown in Fig. 5.

B. On the topic of Analysis 4

The differences between the benchmark times is not significant when compared to multi-threaded work. We identified an average of 4.54% time decrease, which is most likely attributed to how the single-threaded workload had go-routines for Alive Cells

and Key presses, whereas the single worker was Step 1 of Distributed, being "barebones".

VIII. IMPLEMENTING HALO EXCHANGE

TO BE IMPLEMENTED

IX. BENCHMARKING AND CRITICAL ANALYSIS OF THE HALO EXCHANGE

LOrem ipsum dolor sit amet, consectetur adipiscing elit. TO BE IMPLEMENTED

X. BENCHMARKING AND CRITICAL ANALYSIS OF ENABLING SDL

- show graphs and describe shape and differences of with vs without SDL
- explain how SDL works
- explain how this causes bottlenecks

XI. COMPARING THE PARRALEL BENCHMARKS

Maecenas sed ultricies felis. Sed imperdiet dictum arcu a egestas.

- Donec dolor arcu, rutrum id molestie in, viverra sed diam
- Curabitur feugiat
- turpis sed auctor facilisis
- arcu eros accumsan lorem, at posuere mi diam sit amet tortor
- Fusce fermentum, mi sit amet euismod rutrum
- sem lorem molestie diam, iaculis aliquet sapien tortor non nisi
- Pellentesque bibendum pretium aliquet

LOrem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. LOrem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris.

Table 1: Example table

Name		
First name	Last Name	Grade
John	Doe	7.5
Richard	Miles	2

Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Text requiring further explanation¹.

XII. RESULTS

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

$$e = mc^2 \quad (1)$$

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

¹Example footnote

XIII. DISCUSSION

A. Areas for improvement

A statement requiring citation [Figueredo and Wolf, 2009]. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

B. Thanks

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

REFERENCES

- [Figueredo and Wolf, 2009] Figueredo, A. J. and Wolf, P. S. A. (2009). Assortative pairing and life history strategy - a cross-cultural study. *Human Nature*, 20:317–330.